

Assignment 1: Part 2

Name: Bismay Parija

Roll Number: 20CS30067

Import Libraries

```
In [61]: # Import the necessary libraries
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score
```

Data Preprocessing

```
In [62]: # Load the dataset
df = pd.read_csv('../dataset/cross-validation.csv')
print(df.shape)
df.info()
```

```
(614, 13)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Loan_ID               614 non-null   object
 1   Gender                601 non-null   object
 2   Married               611 non-null   object
 3   Dependents            599 non-null   object
 4   Education             614 non-null   object
 5   Self_Employed         582 non-null   object
 6   ApplicantIncome       614 non-null   int64
 7   CoapplicantIncome     614 non-null   float64
 8   LoanAmount            592 non-null   float64
 9   Loan_Amount_Term      600 non-null   float64
10   Credit_History         564 non-null   float64
11   Property_Area         614 non-null   object
12   Loan_Status           614 non-null   object
dtypes: float64(4), int64(1), object(8)
memory usage: 62.5+ KB
```

```
In [63]: # Handle missing values
# Fill missing numerical values with the mean and categorical values with the mode.
df['Gender'].fillna(df['Gender'].mode()[0], inplace=True)
df['Married'].fillna(df['Married'].mode()[0], inplace=True)
df['Dependents'].fillna(df['Dependents'].mode()[0], inplace=True)
```

```
df['Self_Employed'].fillna(df['Self_Employed'].mode()[0], inplace=True)
df['LoanAmount'].fillna(df['LoanAmount'].mean(), inplace=True)
df['Loan_Amount_Term'].fillna(df['Loan_Amount_Term'].mean(), inplace=True)
df['Credit_History'].fillna(df['Credit_History'].mode()[0], inplace=True)
```

```
In [64]: # One-hot encode categorical variables
df = pd.get_dummies(df, columns=['Loan_ID', 'Gender', 'Married', 'Dependents', 'Educ
```

```
In [65]: # Split the dataset into 80% for training and 20% for testing
X = df.drop('Loan_Status', axis=1)
y = df['Loan_Status']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.2, random_stat
# X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=.12
```

Train Logistic Regression Model

```
In [66]: # Train a Logistic Regression model with the saga solver and no regularization
model = LogisticRegression(solver='saga', penalty=None, random_state=42)
model.fit(X_train, y_train)
```

c:\Users\HP\anaconda3\lib\site-packages\sklearn\linear_model_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(

```
Out[66]: LogisticRegression
LogisticRegression(penalty=None, random_state=42, solver='saga')
```

K-Fold Cross Validation

```
In [67]: # Define the number of folds
num_folds = 5
fold_size = len(X_train) // num_folds

# Initialize lists to store metric scores
accuracy_scores = []
precision_scores = []
recall_scores = []

# Perform cross-validation
for i in range(num_folds):
    # Split the data into train and validation sets for this fold
    start_idx = i * fold_size
    end_idx = (i + 1) * fold_size
    val_X = X_train.iloc[start_idx:end_idx]
    val_y = y_train.iloc[start_idx:end_idx]
    train_X = pd.concat([X_train.iloc[:start_idx], X_train.iloc[end_idx:]]
    train_y = pd.concat([y_train.iloc[:start_idx], y_train.iloc[end_idx:]]

    # Train the model on the training set
    model.fit(train_X, train_y)

    # Predict on the validation set
```

```

val_pred = model.predict(val_X)

# Calculate accuracy, precision, and recall for this fold
accuracy = accuracy_score(val_y, val_pred)
precision = precision_score(val_y, val_pred, pos_label='Y')
recall = recall_score(val_y, val_pred, pos_label='Y')

accuracy_scores.append(accuracy)
precision_scores.append(precision)
recall_scores.append(recall)

```

```

c:\Users\HP\anaconda3\lib\site-packages\sklearn\linear_model\_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
  warnings.warn(
c:\Users\HP\anaconda3\lib\site-packages\sklearn\linear_model\_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
  warnings.warn(
c:\Users\HP\anaconda3\lib\site-packages\sklearn\linear_model\_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
  warnings.warn(
c:\Users\HP\anaconda3\lib\site-packages\sklearn\linear_model\_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
  warnings.warn(
c:\Users\HP\anaconda3\lib\site-packages\sklearn\linear_model\_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
  warnings.warn(

```

Calculating Metrics

```

In [68]: # Calculate mean scores across all folds
mean_accuracy = np.mean(accuracy_scores)
mean_precision = np.mean(precision_scores)
mean_recall = np.mean(recall_scores)

# Print the mean scores
print("Mean Accuracy:", mean_accuracy)
print("Mean Precision:", mean_precision)
print("Mean Recall:", mean_recall)

```

```

Mean Accuracy: 0.6959183673469387
Mean Precision: 0.6959183673469387
Mean Recall: 1.0

```